Decentralized curation

Putting self-interest to work for the public good.

Institute of Open Science Practices (IOSP) 2026, Leiden · Wednesday 14 October 2026, 14:00 to 15:20

Informed judgment, public purpose

Credible consensus is evidence-informed judgment under an explicit policy, with correction paths. Reward it to support better funding decisions. It does not certify truth or public benefit; those need evidence of outcomes for affected people.

The full mechanism

1. **Set terms:** funding budget, public policy and fixed economic rules.
2. **Nominate:** identify a beneficiary, publish evidence, request a weight and deposit the funding asset.
3. **Judge relevance:** stake-backed, randomly drawn juries seal and reveal scores under the policy.
4. **Appeal:** crowdfund score proposals during a bounded window. Unopposed: forfeiture, full refund, no fee or profit. Opposed: jury target doubles. Payouts follow inverse distance to the final score.
5. **Fund:** eligible weights split the continuing flow after any required holdback.
6. **Challenge and revisit:** once scored, claims can face external arbitration when no relevance case is open. Re-scoring can change future funding.

Weight is not a funding share

Effective weight = the lower of requested weight and curated score. Score ratios set funding weights, not ratios of public benefit. A sole eligible positive weight takes the full flow. The NumPy example is illustrative. The pool pays jury fees first; contest funds cover shortfalls.

Dependencies and authority

Good outcomes need a defensible policy, informed jurors, usable evidence and affordable challenges. Stake does not certify expertise; agreement does not establish public benefit.

Where configured, a temporary circuit-breaker holder may permanently stop future accrual once, before expiry. They cannot take prior accrual.

Try, read, stay in touch

The room

Published preprint

Mailing list

Today: experience one layer

Read the policy, score alone, seal, then reveal. This hosted game uses equal seats and points. Its band follows the mean and standard deviation; rewards depend on closeness. At essentially zero spread, distance forfeits switch off.

Workshop simplification: each round with opened scores pays the full base reward, even when those scores all agree.

Use an alias or observe. Revealed scores and standings are visible; adversarial role-play may be assigned. Ask before photographing results. The host wipes session results afterward; they have no automatic expiry.

Keep score proofs private. This game's reporting bounty is absent from the inspected contracts. The game does not run accuracy disputes, jury selection, appeals or a funding stream.

Example pool policy

How much does this nominee's completed and maintained work underpin how research is done, shared and kept, now?

Score the work identified by each nomination. If entries claim the same work, credit that shared work once across them; distinct contributions by the same beneficiary remain eligible.

Docket: Inkscape, Linux, PDF, LaTeX, arXiv, NumPy.

From experience to evidence

Ask what you judged, what the game rewarded and what would justify another jury. This exercise does not establish better decisions or resistance to collusion.

Arrow is a planned precursor experiment at Octant: invited experts score a real docket under a published policy. It assumes honesty, competence and non-collusion; it does not validate the full permissionless system. octant.build

Room: decentralized-curation.0xferit.workers.dev/r/leiden

Join: decentralized-curation.0xferit.workers.dev/join/

Published preprint, version 1.2.0: doi.org/10.5281/zenodo.23313625.

Release source · Public mechanism and evidence limits. Later hardening is outside this archived release.

Octant

Sponsored by Octant.

Decentralized Curation research by Ferit Tunçer.

Copyright (C) 2026 Ferit Tunçer. Licence: code GPL-3.0-or-later, original text and figures CC BY 4.0; third-party marks excluded. https://decentralized-curation.0xferit.workers.dev