WORKSHOP HANDOUT

Decentralized curation

Putting self-interest to work for the public good.

Institute of Open Science Practices (IOSP) 2026, Leiden · Wednesday 14 October 2026, 14:00 to 15:20

Who judges what matters?

Research depends on judgments about what to read, cite, reuse and reward. This proposal uses private incentives for assessment and correction to pursue credible consensus. Public-goods funding is one application.

The funding application

1. **Set terms:** funding budget, public policy and fixed economic rules.
2. **Nominate:** identify a beneficiary, publish evidence, request a weight and post a deposit. Claims need no prior truth approval.
3. **Judge relevance:** stake-backed, randomly drawn juries seal and reveal scores under the policy.
4. **Appeal:** publicly crowdfund competing scores. An opposed appeal doubles the jury target; payouts depend on distance to the final score.
5. **Fund:** eligible weights split the continuing flow after any required holdback.
6. **Fact-check:** after scoring, anyone can post a bond and fees to contest false or non-falsifiable claims, if no relevance case is open. An external arbitrator decides. A successful challenger takes the nomination's bond and ends its pool benefits. The author may nominate again.
7. **Revisit:** new relevance assessments can change future funding.

Weight is not a funding share

Effective weight = the lower of requested weight and curated score. These are funding ratios, not measured ratios of public benefit. A sole eligible positive weight takes the full flow. Jury fees come first; contest funds cover shortfalls. See the illustrative NumPy example.

A configured, temporary circuit-breaker holder can permanently stop future accrual once, before expiry, but cannot take prior accrual.

Today: experience one layer

Read the policy, score alone, seal, then reveal. This hosted game uses equal seats and points. Its band follows the mean and standard deviation; rewards depend on closeness. At essentially zero spread, distance forfeits switch off.

Workshop simplification: each round with opened scores pays the full base reward, even when those scores all agree.

Use an alias or observe. Revealed scores and standings are visible; adversarial role-play may be assigned. Ask before photographing results. The host wipes session results afterward; they have no automatic expiry.

Keep score proofs private. This game's reporting bounty is absent from the inspected contracts. The game does not run accuracy disputes, jury selection, appeals or a funding stream.

Example pool policy

How much does this nominee's completed and maintained work underpin how research is done, shared and kept, now?

Score the work identified by each nomination. If entries claim the same work, credit that shared work once across them; distinct contributions by the same beneficiary remain eligible.

Docket: Inkscape, Linux, PDF, LaTeX, arXiv, NumPy.

From experience to evidence

Ask whether the game rewarded evidence use or merely agreement. Neither stake nor agreement proves truth or public benefit. Reliable judgments and collusion resistance remain unproven.

Arrow is a planned precursor experiment at Octant: invited experts score a real docket under a published policy. It assumes honesty, competence and non-collusion; it does not validate the full permissionless system. octant.build

Try, read, stay in touch

The room

Published preprint

Mailing list

Room: decentralized-curation.0xferit.workers.dev/r/leiden

Join: decentralized-curation.0xferit.workers.dev/join/

Published preprint, version 1.2.0: doi.org/10.5281/zenodo.23313625.

Release source · Public mechanism and evidence limits. Later hardening is outside this archived release.

Sponsored by Octant

Decentralized Curation research by Ferit Tunçer.

Copyright (C) 2026 Ferit Tunçer. Licence: code GPL-3.0-or-later, original text and figures CC BY 4.0; third-party marks excluded. https://decentralized-curation.0xferit.workers.dev